

```
!pip install pandas
!pip install numpy
!pip install matplotlib
!pip install seaborn
!pip install sklearn
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.2.2)
Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/dist-packages (from pandas) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.26.4)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.8.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.54.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.7)
Requirement already satisfied: numpy<2,>=1.21 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.26.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.13.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in /usr/local/lib/python3.10/dist-packages (from seaborn) (1.26.4)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.10/dist-packages (from seaborn) (2.2.2)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /usr/local/lib/python3.10/dist-packages (from seaborn) (3.8.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.54.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.7)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn) (2024.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)
Collecting sklearn
  Using cached sklearn-0.0.post12.tar.gz (2.6 kB)
  error: subprocess-exited-with-error

  × python setup.py egg_info did not run successfully.
  | exit code: 1
  | See above for output.

  note: This error originates from a subprocess, and is likely not a problem with pip.
  Preparing metadata (setup.py) ... error
error: metadata-generation-failed

× Encountered error while generating package metadata.
  See above for output.

note: This is an issue with the package mentioned above, not pip.
hint: See above for details.
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from google.colab import files
uploaded = files.upload()
in_vehicle_coupon_data = pd.read_csv("in-vehicle-coupon-recommendation.csv")
in_vehicle_coupon_data.head() #Display the first few rows
```

Choose Files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving in-vehicle-coupon-recommendation.csv to in-vehicle-coupon-recommendation (2).csv

	destination	passanger	weather	temperature	time	coupon	expiration	gender	age	maritalStatus	...	CoffeeHouse	CarryAway	RestaurantLe
0	No Urgent Place	Alone	Sunny	55	2PM	Restaurant(<20)	1d	Female	21	Unmarried partner	...	never	NaN	
1	No Urgent Place	Friend(s)	Sunny	80	10AM	Coffee House	2h	Female	21	Unmarried partner	...	never	NaN	
2	No Urgent Place	Friend(s)	Sunny	80	10AM	Carry out & Take away	2h	Female	21	Unmarried partner	...	never	NaN	
3	No Urgent Place	Friend(s)	Sunny	80	2PM	Coffee House	2h	Female	21	Unmarried partner	...	never	NaN	
4	No Urgent Place	Friend(s)	Sunny	80	2PM	Coffee House	1d	Female	21	Unmarried partner	...	never	NaN	

5 rows × 26 columns

in_vehicle_coupon_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12684 entries, 0 to 12683
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   destination            12684 non-null  object
1   passanger              12684 non-null  object
2   weather                12684 non-null  object
3   temperature            12684 non-null  int64
4   time                   12684 non-null  object
5   coupon                 12684 non-null  object
6   expiration             12684 non-null  object
7   gender                 12684 non-null  object
8   age                    12684 non-null  object
9   maritalStatus          12684 non-null  object
10  has_children           12684 non-null  int64
11  education              12684 non-null  object
12  occupation              12684 non-null  object
13  income                 12684 non-null  object
14  car                    108 non-null    object
15  Bar                    12577 non-null  object
16  CoffeeHouse            12467 non-null  object
17  CarryAway              12533 non-null  object
18  RestaurantLessThan20   12554 non-null  object
19  Restaurant20To50       12495 non-null  object
20  toCoupon_GEQ5min       12684 non-null  int64
21  toCoupon_GEQ15min      12684 non-null  int64
22  toCoupon_GEQ25min      12684 non-null  int64
23  direction_same         12684 non-null  int64
24  direction_opp          12684 non-null  int64
25  Y                      12684 non-null  int64
dtypes: int64(8), object(18)
memory usage: 2.5+ MB
```

#This code snippet has been taken from Niralidedaniya (2022).

```
print('Is there any missing value present or not?',in_vehicle_coupon_data.isnull().values.any())
missing_percentage = in_vehicle_coupon_data.isnull().sum()*100/len(in_vehicle_coupon_data) #Calculate the percentage of missing data
missing_value_df = pd.DataFrame({'missing_count': in_vehicle_coupon_data.isnull().sum(),'missing_percentage': missing_percentage}) #create a dataframe
print(missing_value_df) #print the dataframe created in the previous step
```

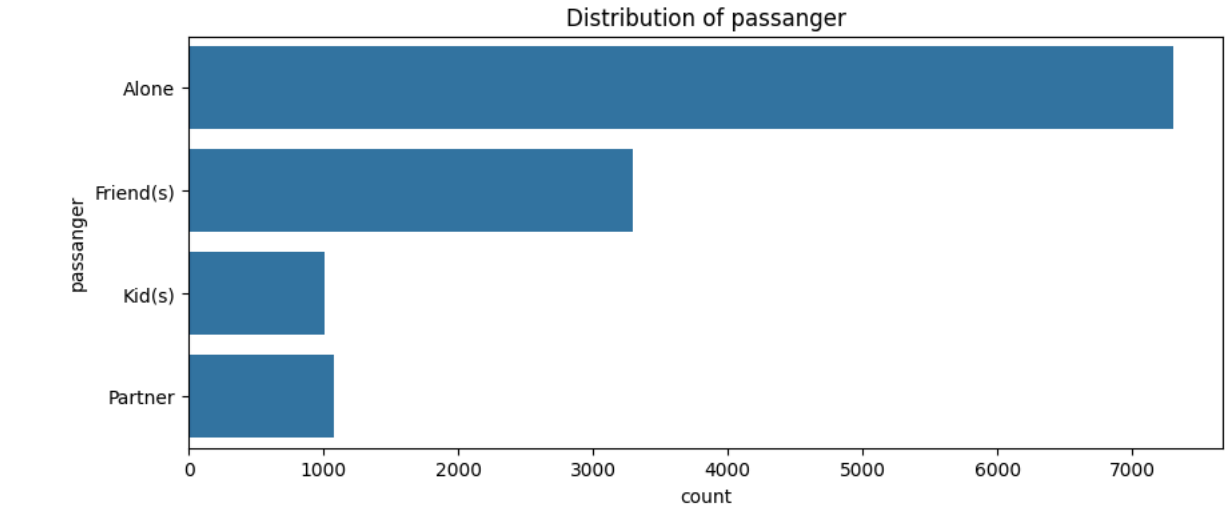
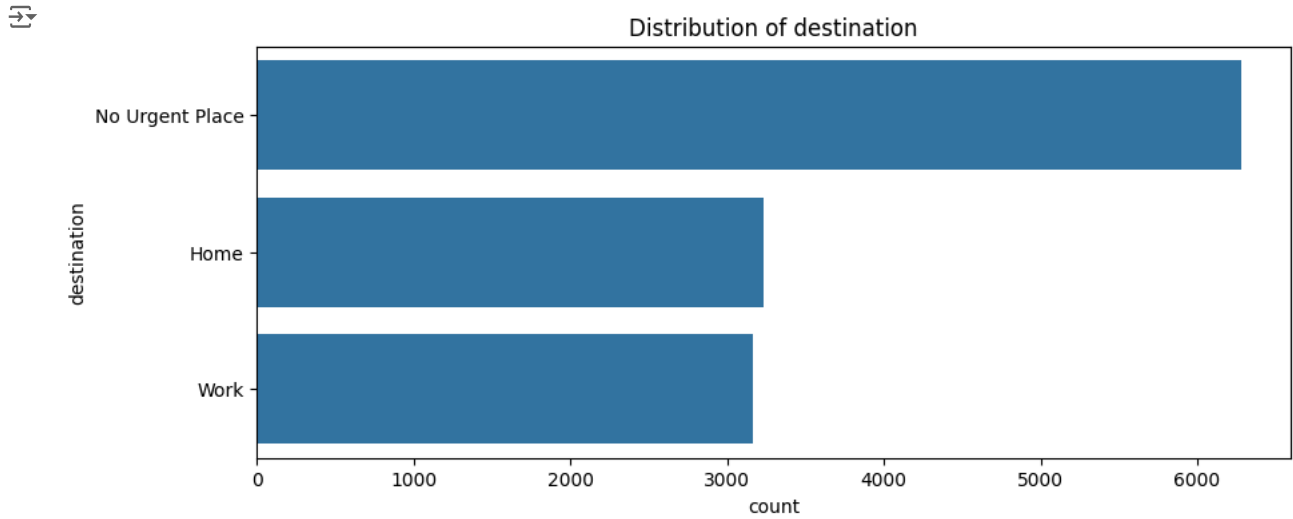
```
Is there any missing value present or not? True
missing_count missing_percentage
destination      0      0.000000
passanger        0      0.000000
weather          0      0.000000
temperature      0      0.000000
time             0      0.000000
coupon           0      0.000000
expiration       0      0.000000
gender           0      0.000000
age              0      0.000000
maritalStatus    0      0.000000
has_children     0      0.000000
education        0      0.000000
occupation       0      0.000000
income           0      0.000000
car              12576   99.148534
Bar              107     0.843582
CoffeeHouse      217     1.710817
CarryAway        151     1.190476
```

RestaurantLessThan20	130	1.024913
Restaurant20To50	189	1.490066
toCoupon_GEQ5min	0	0.000000
toCoupon_GEQ15min	0	0.000000
toCoupon_GEQ25min	0	0.000000
direction_same	0	0.000000
direction_opp	0	0.000000
Y	0	0.000000


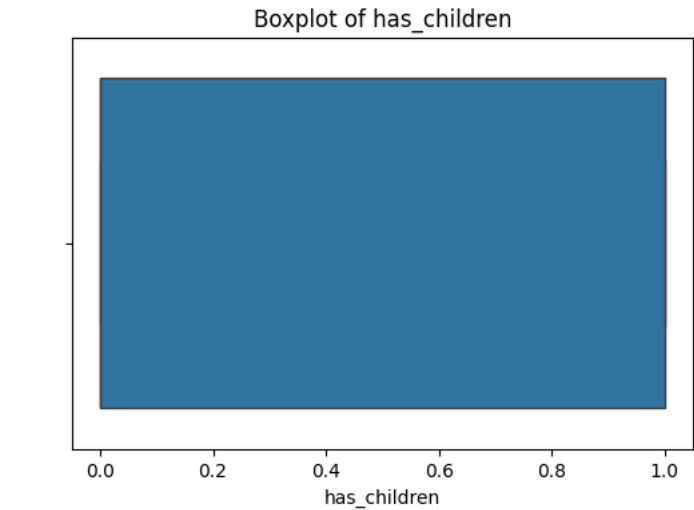

```
#This code snippet has been taken from Niralidedaniya (2022).
#For the target column, calculate the percentage of coupons accepted and rejected.
Y_counts = in_vehicle_coupon_data.groupby('Y').Y.count()
print('Accepted coupons:',Y_counts[1],round(Y_counts[1]/in_vehicle_coupon_data.shape[0]*100,3),'%')
print('Rejected coupons:',Y_counts[0],round(Y_counts[0]/in_vehicle_coupon_data.shape[0]*100,3),'%')
```

Accepted coupons: 7210 56.843 %
Rejected coupons: 5474 43.157 %

```
#This code snippet came from Inyama (2023).
#We want to see the distribution of data in each of the columns that have categorical data
categorical_variables = in_vehicle_coupon_data.select_dtypes(include=['object']).columns
for var in categorical_variables:
    plt.figure(figsize=(10, 4))
    sns.countplot(y=var, data=in_vehicle_coupon_data)
    plt.title(f'Distribution of {var}')
    plt.show()
```



```
#This code snippet came from Inyama (2023).
#We want to see the box plot for numerical data in numerical columns
numerical_columns = in_vehicle_coupon_data.select_dtypes(include=['int64', 'float64']).columns.tolist()
numerical_columns.remove('Y')
for column in numerical_columns:
    plt.figure(figsize=(6, 4))
    sns.boxplot(x=in_vehicle_coupon_data[column])
    plt.title(f'Boxplot of {column}')
    plt.show()
```



```
#Drop column car as it has too many missing values
#Drop column toCoupon_GEQ5min because there is no variability in its value
#Drop column occupation because it has too many categories that leads to a lot of noise
in_vehicle_coupon_data = in_vehicle_coupon_data.drop(['car', 'occupation', 'toCoupon_GEQ5min'], axis=1)
in_vehicle_coupon_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12684 entries, 0 to 12683
Data columns (total 23 columns):
#   Column          Non-Null Count  Dtype
---  -
0   destination      12684 non-null  object
1   passanger        12684 non-null  object
2   weather          12684 non-null  object
3   temperature      12684 non-null  int64
4   time             12684 non-null  object
5   coupon           12684 non-null  object
```

```
6 expiration 12684 non-null object
7 gender      12684 non-null object
8 age         12684 non-null object
9 maritalStatus 12684 non-null object
10 has_children 12684 non-null int64
11 education   12684 non-null object
12 income      12684 non-null object
13 Bar         12577 non-null object
14 CoffeeHouse 12467 non-null object
15 CarryAway   12533 non-null object
16 RestaurantLessThan20 12554 non-null object
17 Restaurant20To50 12495 non-null object
18 toCoupon_GEQ15min 12684 non-null int64
19 toCoupon_GEQ25min 12684 non-null int64
20 direction_same 12684 non-null int64
21 direction_opp 12684 non-null int64
22 Y           12684 non-null int64
```

```
dtypes: int64(7), object(16)
memory usage: 2.2+ MB
```

```
# frequent value / mode imputation for missing values in data. This code snippet has been taken from Niralidedaniya (2022).
in_vehicle_coupon_data['Bar']=in_vehicle_coupon_data['Bar'].fillna(in_vehicle_coupon_data['Bar'].value_counts().index[0])
in_vehicle_coupon_data['CoffeeHouse']=in_vehicle_coupon_data['CoffeeHouse'].fillna(in_vehicle_coupon_data['CoffeeHouse'].value_counts().index[0])
in_vehicle_coupon_data['CarryAway']=in_vehicle_coupon_data['CarryAway'].fillna(in_vehicle_coupon_data['CarryAway'].value_counts().index[0])
in_vehicle_coupon_data['RestaurantLessThan20']=in_vehicle_coupon_data['RestaurantLessThan20'].fillna(in_vehicle_coupon_data['RestaurantLessThan20'].value_
in_vehicle_coupon_data['Restaurant20To50']=in_vehicle_coupon_data['Restaurant20To50'].fillna(in_vehicle_coupon_data['Restaurant20To50'].value_counts().inde
#Lets check for missing values again
print('Is there any missing value present?',in_vehicle_coupon_data.isnull().values.any())
```

Is there any missing value present? False

```
from sklearn.preprocessing import OrdinalEncoder
encoder = OrdinalEncoder()
encoded_in_vehicle_coupon_data = encoder.fit_transform(in_vehicle_coupon_data)
encoded_in_vehicle_coupon_data = pd.DataFrame(encoded_in_vehicle_coupon_data, columns=in_vehicle_coupon_data.columns)
encoded_in_vehicle_coupon_data.cov() #covariance
```

	destination	passanger	weather	temperature	time	coupon	expiration	gender	age	maritalStatus	...	Bar
destination	0.504658	-0.078836	-0.035245	-0.015041	0.420743	-0.010986	-0.011497	0.002407	-0.003339	0.001797	...	-0.005564
passanger	-0.078836	0.887317	0.035703	0.040508	-0.543328	0.023840	0.034968	-0.009838	0.010415	-0.016771	...	0.022005
weather	-0.035245	0.035703	0.401445	0.210926	-0.021266	0.125609	0.005569	-0.008552	-0.027806	-0.008874	...	0.010249
temperature	-0.015041	0.040508	0.210926	0.587031	-0.065117	0.133050	0.047207	-0.009767	-0.046974	0.002288	...	0.010868
time	0.420743	-0.543328	-0.021266	-0.065117	2.072183	0.092426	-0.060532	-0.002523	-0.025719	0.005786	...	-0.009895
coupon	-0.010986	0.023840	0.125609	0.133050	0.092426	1.818577	0.099355	0.004809	0.006944	-0.000149	...	-0.015862
expiration	-0.011497	0.034968	0.005569	0.047207	-0.060532	0.099355	0.246532	-0.000314	0.007207	-0.005098	...	-0.006669
gender	0.002407	-0.009838	-0.008552	-0.009767	-0.002523	0.004809	-0.000314	0.249842	-0.067526	0.023960	...	-0.120201
age	-0.003339	0.010415	-0.027806	-0.046974	-0.025719	0.006944	0.007207	-0.067526	4.950184	-0.328927	...	0.712399
maritalStatus	0.001797	-0.016771	-0.008874	0.002288	0.005786	-0.000149	-0.005098	0.023960	-0.328927	0.693751	...	-0.118302
has_children	-0.002347	0.016028	0.003950	-0.007441	-0.005145	-0.006923	0.003918	-0.039384	0.335703	-0.177913	...	0.127863
education	0.011758	0.001656	0.015185	0.022296	-0.008154	-0.001837	-0.008597	0.014140	0.363653	0.107978	...	0.116209
income	-0.025464	-0.005402	-0.051117	-0.047347	-0.022016	0.002894	-0.013671	0.032841	0.267308	0.150888	...	0.053494
Bar	-0.005564	0.022005	0.010249	0.010868	-0.009895	-0.015862	-0.006669	-0.120201	0.712399	-0.118302	...	2.407485
CoffeeHouse	-0.007850	-0.012488	-0.005865	0.008828	-0.011495	-0.004406	-0.009675	0.045669	-0.002177	0.014938	...	0.378947
CarryAway	-0.005959	-0.017730	-0.026723	-0.027345	0.001266	-0.004440	-0.002079	-0.009576	0.216170	-0.022649	...	0.134812
RestaurantLessThan20	0.002886	-0.033238	-0.003905	-0.001147	-0.004725	0.016685	-0.006039	0.025848	-0.074016	-0.001212	...	0.139119
Restaurant20To50	-0.000269	-0.042455	0.005444	0.002082	0.019281	0.011179	-0.000336	-0.000155	-0.004063	0.055322	...	0.375264
toCoupon_GEQ15min	0.049593	0.030170	-0.038262	-0.059057	0.005156	-0.088045	0.010531	-0.001743	0.029335	-0.020447	...	0.015285
toCoupon_GEQ25min	0.045391	-0.060297	-0.041579	-0.053675	0.136276	-0.049269	-0.005304	0.000444	-0.000045	0.001348	...	0.002514
direction_same	-0.024310	-0.103995	0.004609	0.030548	0.184128	-0.040432	0.006848	-0.000923	-0.007565	0.005645	...	-0.000763
direction_opp	0.024310	0.103995	-0.004609	-0.030548	-0.184128	0.040432	-0.006848	0.000923	0.007565	-0.005645	...	0.000763
Y	-0.000671	0.024082	0.031006	0.023241	-0.033780	0.064804	-0.031952	0.010886	-0.038836	0.010348	...	-0.058434

23 rows × 23 columns

Since direction_same has same covariance values as direction_opp. It makes sense to just have one on them and reduce the noise.

```
in_vehicle_coupon_data = in_vehicle_coupon_data.drop(['direction_opp'], axis=1)
```

```
in_vehicle_coupon_data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12684 entries, 0 to 12683
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   destination                           12684 non-null  object
1   passanger                             12684 non-null  object
2   weather                               12684 non-null  object
3   temperature                           12684 non-null  int64
4   time                                  12684 non-null  object
5   coupon                                12684 non-null  object
6   expiration                            12684 non-null  object
7   gender                                12684 non-null  object
8   age                                   12684 non-null  object
9   maritalStatus                         12684 non-null  object
10  has_children                          12684 non-null  int64
11  education                             12684 non-null  object
12  income                                12684 non-null  object
13  Bar                                    12684 non-null  object
14  CoffeeHouse                           12684 non-null  object
15  CarryAway                             12684 non-null  object
16  RestaurantLessThan20                  12684 non-null  object
17  Restaurant20To50                      12684 non-null  object
18  toCoupon_GEQ15min                     12684 non-null  int64
19  toCoupon_GEQ25min                     12684 non-null  int64
20  direction_same                        12684 non-null  int64
21  Y                                      12684 non-null  int64
dtypes: int64(6), object(16)
memory usage: 2.1+ MB
```

0 250 500 750 1000 1250 1500 1750 2000

#Get the correlation matrix using ordinal encoding for the categorical values

```
encoder = OrdinalEncoder()
```

```
encoded_in_vehicle_coupon_data = encoder.fit_transform(in_vehicle_coupon_data)
```

```
encoded_in_vehicle_coupon_data = pd.DataFrame(encoded_in_vehicle_coupon_data, columns=in_vehicle_coupon_data.columns)
```

```
correlation_matrix = encoded_in_vehicle_coupon_data.corr()
```

```
print(correlation_matrix)
```



```
destination    destination    passanger    weather    temperature    time \
passanger      -0.117811    1.000000    0.059821    0.056127    -0.400690
weather        -0.078305    0.059821    1.000000    0.434497    -0.023316
temperature    -0.027633    0.056127    0.434497    1.000000    -0.059041
time           0.411437    -0.400690    -0.023316    -0.059041    1.000000
coupon         -0.011468    0.018767    0.147008    0.128771    0.047612
expiration     -0.032594    0.074764    0.017702    0.124090    -0.084691
gender         0.006779    -0.020896    -0.027003    -0.025504    -0.003507
age            -0.002112    0.004969    -0.019725    -0.027556    -0.008030
maritalStatus  0.003036    -0.021376    -0.016816    0.003585    0.004826
has_children   -0.006707    0.034542    0.012657    -0.019716    -0.007256
education      0.008793    0.000934    0.012733    0.015460    -0.003009
income         -0.014554    -0.002329    -0.032758    -0.025091    -0.006210
Bar            -0.005048    0.015055    0.010426    0.009142    -0.004430
CoffeeHouse    -0.007185    -0.008619    -0.006018    0.007491    -0.005192
CarryAway      -0.007676    -0.017223    -0.038593    -0.032657    0.000804
RestaurantLessThan20  0.003497    -0.030378    -0.005306    -0.001289    -0.002826
Restaurant20To50 -0.000255    -0.030401    0.005795    0.001833    0.009035
toCoupon_GEQ15min  0.140684    0.064544    -0.121698    -0.155332    0.007218
toCoupon_GEQ25min  0.197240    -0.197595    -0.202572    -0.216254    0.292231
direction_same  -0.083328    -0.268830    0.017712    0.097085    0.311467
Y              -0.001906    0.051614    0.098800    0.061240    -0.047377

coupon    expiration    gender    age    maritalStatus \
destination    -0.011468    -0.032594    0.006779    -0.002112    0.003036
passanger      0.018767    0.074764    -0.020896    0.004969    -0.021376
weather        0.147008    0.017702    -0.027003    -0.019725    -0.016816
temperature    0.128771    0.124090    -0.025504    -0.027556    0.003585
time           0.047612    -0.084691    -0.003507    -0.008030    0.004826
coupon         1.000000    0.148383    0.007134    0.002314    -0.000132
expiration     0.148383    1.000000    -0.001264    0.006523    -0.012328
gender         0.007134    -0.001264    1.000000    -0.060720    0.057552
age            0.002314    0.006523    -0.060720    1.000000    -0.177495
maritalStatus  -0.000132    -0.012328    0.057552    -0.177495    1.000000
has_children   -0.010422    0.016020    -0.159956    0.306306    -0.433628
education      -0.000724    -0.009198    0.015029    0.086833    0.068872
income         0.000871    -0.011180    0.026677    0.048782    0.073555
Bar            -0.007581    -0.008656    -0.154986    0.206363    -0.091539
CoffeeHouse    -0.002124    -0.012669    0.059404    -0.000636    0.011661
CarryAway      -0.003012    -0.003832    -0.017530    0.088903    -0.024881
RestaurantLessThan20  0.010651    -0.010471    0.044519    -0.028640    -0.001253
Restaurant20To50  0.005592    -0.000457    -0.000209    -0.001232    0.044802
toCoupon_GEQ15min -0.131571    0.042740    -0.007028    0.026571    -0.049471
```



```

destination      0.0569      0.007      8.321      0.000      0.044      0.070
passanger        0.0471      0.005      9.367      0.000      0.037      0.057
weather          0.1170      0.007     16.382      0.000      0.103      0.131
temperature      0.0266      0.006      4.160      0.000      0.014      0.039
time            -0.0096      0.004     -2.434      0.015     -0.017     -0.002
coupon           0.0541      0.003     16.362      0.000      0.048      0.061
expiration       -0.1407      0.009    -15.771      0.000     -0.158     -0.123
gender           0.0866      0.009     10.012      0.000      0.070      0.104
age              0.0025      0.002      1.189      0.234     -0.002      0.007
maritalStatus    0.0635      0.005     12.183      0.000      0.053      0.074
has_children     0.0444      0.010      4.483      0.000      0.025      0.064
education        0.0200      0.002      8.762      0.000      0.016      0.025
income           0.0014      0.002      0.820      0.412     -0.002      0.005
Bar              0.0008      0.003      0.281      0.779     -0.005      0.007
CoffeeHouse     -0.0363      0.003    -12.574      0.000     -0.042     -0.031
CarryAway        0.0014      0.004      0.354      0.723     -0.006      0.009
RestaurantLessThan20 0.0150      0.004      3.935      0.000      0.008      0.022
Restaurant20To50 -0.0019      0.003     -0.638      0.524     -0.008      0.004
toCoupon_GEQ15min 0.0237      0.009      2.502      0.012      0.005      0.042
toCoupon_GEQ25min -0.0321      0.016     -2.039      0.041     -0.063     -0.001
direction_same   0.1026      0.013      8.132      0.000      0.078      0.127
=====
Omnibus:          77704.989   Durbin-Watson:      1.682
Prob(Omnibus):    0.000   Jarque-Bera (JB):   1296.009
Skew:             -0.228   Prob(JB):           3.76e-282
Kurtosis:         1.502   Cond. No.           31.5
=====

```

Notes:
[1] R² is computed without centering (uncentered) since the model does not contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

I choose alpha to be 0.05. Coefficients having a p-value of 0.05 or less will be statistically significant. Bar has 0.779 for p-value. Bar will be dropped. It is the least statistically significant.

```

x_stepwise_2 = x_stepwise.drop("Bar", axis=1)
result_stepwise_2 = sm.OLS(y_stepwise, x_stepwise_2).fit()
print(result_stepwise_2.summary())

```

```

OLS Regression Results

=====
Dep. Variable:      Y      R-squared (uncentered):      0.587
Model:              OLS      Adj. R-squared (uncentered):      0.586
Method:              Least Squares      F-statistic:      899.5
Date:                Mon, 11 Nov 2024      Prob (F-statistic):      0.00
Time:                05:49:27      Log-Likelihood:      -8808.9
No. Observations:    12684      AIC:      1.766e+04
Df Residuals:        12664      BIC:      1.781e+04
Df Model:            20
Covariance Type:     nonrobust
=====
                    coef      std err      t      P>|t|      [0.025      0.975]
-----
destination      0.0570      0.007      8.332      0.000      0.044      0.070
passanger        0.0472      0.005      9.400      0.000      0.037      0.057
weather          0.1172      0.007     16.463      0.000      0.103      0.131
temperature      0.0266      0.006      4.172      0.000      0.014      0.039
time            -0.0095      0.004     -2.426      0.015     -0.017     -0.002
coupon           0.0541      0.003     16.380      0.000      0.048      0.061
expiration       -0.1407      0.009    -15.770      0.000     -0.158     -0.123
gender           0.0863      0.009     10.041      0.000      0.069      0.103
age              0.0026      0.002      1.265      0.206     -0.001      0.007
maritalStatus    0.0636      0.005     12.229      0.000      0.053      0.074
has_children     0.0448      0.010      4.567      0.000      0.026      0.064
education        0.0201      0.002      8.790      0.000      0.016      0.025
income           0.0014      0.002      0.833      0.405     -0.002      0.005
CoffeeHouse     -0.0362      0.003    -12.713      0.000     -0.042     -0.031
CarryAway        0.0015      0.004      0.368      0.713     -0.006      0.009
RestaurantLessThan20 0.0150      0.004      3.966      0.000      0.008      0.022
Restaurant20To50 -0.0018      0.003     -0.599      0.549     -0.008      0.004
toCoupon_GEQ15min 0.0238      0.009      2.519      0.012      0.005      0.042
toCoupon_GEQ25min -0.0319      0.016     -2.032      0.042     -0.063     -0.001
direction_same   0.1027      0.013      8.147      0.000      0.078      0.127
=====
Omnibus:          77832.632   Durbin-Watson:      1.682
Prob(Omnibus):    0.000   Jarque-Bera (JB):   1295.135
Skew:             -0.228   Prob(JB):           5.82e-282
Kurtosis:         1.503   Cond. No.           29.7
=====

```

Notes:
[1] R² is computed without centering (uncentered) since the model does not contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In the second iteration, CarryAway is least statistically significant because it has highest p-value of 0.713 which is greater than alpha of 0.05. CarryAway should be dropped.

```
x_stepwise_3 = x_stepwise_2.drop("CarryAway", axis=1)
result_stepwise_3 = sm.OLS(y_stepwise, x_stepwise_3).fit()
print(result_stepwise_3.summary())
```

OLS Regression Results

Dep. Variable:	Y	R-squared (uncentered):	0.587
Model:	OLS	Adj. R-squared (uncentered):	0.586
Method:	Least Squares	F-statistic:	946.9
Date:	Mon, 11 Nov 2024	Prob (F-statistic):	0.00
Time:	05:49:27	Log-Likelihood:	-8808.9
No. Observations:	12684	AIC:	1.766e+04
Df Residuals:	12665	BIC:	1.780e+04
Df Model:	19		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
destination	0.0570	0.007	8.337	0.000	0.044	0.070
passanger	0.0472	0.005	9.409	0.000	0.037	0.057
weather	0.1173	0.007	16.477	0.000	0.103	0.131
temperature	0.0266	0.006	4.171	0.000	0.014	0.039
time	-0.0095	0.004	-2.420	0.016	-0.017	-0.002
coupon	0.0541	0.003	16.396	0.000	0.048	0.061
expiration	-0.1406	0.009	-15.767	0.000	-0.158	-0.123
gender	0.0863	0.009	10.040	0.000	0.069	0.103
age	0.0027	0.002	1.312	0.189	-0.001	0.007
maritalStatus	0.0637	0.005	12.246	0.000	0.053	0.074
has_children	0.0448	0.010	4.570	0.000	0.026	0.064
education	0.0201	0.002	8.805	0.000	0.016	0.025
income	0.0015	0.002	0.859	0.390	-0.002	0.005
CoffeeHouse	-0.0361	0.003	-12.757	0.000	-0.042	-0.031
RestaurantLessThan20	0.0151	0.004	4.005	0.000	0.008	0.023
Restaurant20To50	-0.0016	0.003	-0.559	0.576	-0.007	0.004
toCoupon_GEQ15min	0.0240	0.009	2.532	0.011	0.005	0.042
toCoupon_GEQ25min	-0.0319	0.016	-2.030	0.042	-0.063	-0.001
direction_same	0.1029	0.013	8.158	0.000	0.078	0.128

Omnibus:	77965.192	Durbin-Watson:	1.682
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1294.195
Skew:	-0.228	Prob(JB):	9.31e-282
Kurtosis:	1.503	Cond. No.	29.4

Notes:
[1] R² is computed without centering (uncentered) since the model does not contain a constant.
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In the third iteration, Restaurant20To50 is least statistically significant because it has highest p-value of 0.576 which is greater than alpha of 0.05. Restaurant20To50 should be dropped.

```
x_stepwise_4 = x_stepwise_3.drop("Restaurant20To50", axis=1)
result_stepwise_4 = sm.OLS(y_stepwise, x_stepwise_4).fit()
print(result_stepwise_4.summary())
```

OLS Regression Results

Dep. Variable:	Y	R-squared (uncentered):	0.587
Model:	OLS	Adj. R-squared (uncentered):	0.586
Method:	Least Squares	F-statistic:	999.6
Date:	Mon, 11 Nov 2024	Prob (F-statistic):	0.00
Time:	05:49:28	Log-Likelihood:	-8809.1
No. Observations:	12684	AIC:	1.765e+04
Df Residuals:	12666	BIC:	1.779e+04
Df Model:	18		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
destination	0.0569	0.007	8.326	0.000	0.044	0.070
passanger	0.0472	0.005	9.402	0.000	0.037	0.057
weather	0.1169	0.007	16.487	0.000	0.103	0.131
temperature	0.0265	0.006	4.161	0.000	0.014	0.039
time	-0.0096	0.004	-2.436	0.015	-0.017	-0.002
coupon	0.0541	0.003	16.388	0.000	0.048	0.061
expiration	-0.1407	0.009	-15.782	0.000	-0.158	-0.123
gender	0.0862	0.009	10.031	0.000	0.069	0.103
age	0.0026	0.002	1.300	0.193	-0.001	0.007
maritalStatus	0.0633	0.005	12.288	0.000	0.053	0.073

has_children	0.0444	0.010	4.539	0.000	0.025	0.064
education	0.0200	0.002	8.798	0.000	0.016	0.024
income	0.0014	0.002	0.832	0.405	-0.002	0.005
CoffeeHouse	-0.0363	0.003	-13.018	0.000	-0.042	-0.031
RestaurantLessThan20	0.0149	0.004	3.967	0.000	0.008	0.022
toCoupon_GEQ15min	0.0237	0.009	2.510	0.012	0.005	0.042
toCoupon_GEQ25min	-0.0321	0.016	-2.042	0.041	-0.063	-0.001
direction_same	0.1027	0.013	8.146	0.000	0.078	0.127

```
=====
Omnibus:          77733.204   Durbin-Watson:          1.682
Prob(Omnibus):    0.000   Jarque-Bera (JB):          1296.124
Skew:             -0.229   Prob(JB):          3.55e-282
Kurtosis:         1.502   Cond. No.          28.1
=====
```

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
 [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In the fourth iteration, income is least statistically significant because it has highest p-value of 0.405 which is greater than alpha of 0.05. income should be dropped.

```
x_stepwise_5 = x_stepwise_4.drop("income", axis=1)
result_stepwise_5 = sm.OLS(y_stepwise, x_stepwise_5).fit()
print(result_stepwise_5.summary())
```



OLS Regression Results

```
=====
Dep. Variable:          Y   R-squared (uncentered):          0.587
Model:                  OLS   Adj. R-squared (uncentered):          0.586
Method:                 Least Squares   F-statistic:          1058.
Date:                   Mon, 11 Nov 2024   Prob (F-statistic):          0.00
Time:                   05:49:28   Log-Likelihood:          -8809.4
No. Observations:       12684   AIC:          1.765e+04
Df Residuals:           12667   BIC:          1.778e+04
Df Model:                17
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
destination	0.0570	0.007	8.345	0.000	0.044	0.070
passanger	0.0474	0.005	9.456	0.000	0.038	0.057
weather	0.1173	0.007	16.564	0.000	0.103	0.131
temperature	0.0266	0.006	4.172	0.000	0.014	0.039
time	-0.0095	0.004	-2.419	0.016	-0.017	-0.002
coupon	0.0542	0.003	16.455	0.000	0.048	0.061
expiration	-0.1407	0.009	-15.779	0.000	-0.158	-0.123
gender	0.0867	0.009	10.112	0.000	0.070	0.103
age	0.0028	0.002	1.384	0.166	-0.001	0.007
maritalStatus	0.0641	0.005	12.716	0.000	0.054	0.074
has_children	0.0451	0.010	4.632	0.000	0.026	0.064
education	0.0200	0.002	8.825	0.000	0.016	0.024
CoffeeHouse	-0.0364	0.003	-13.027	0.000	-0.042	-0.031
RestaurantLessThan20	0.0152	0.004	4.061	0.000	0.008	0.023
toCoupon_GEQ15min	0.0242	0.009	2.565	0.010	0.006	0.043
toCoupon_GEQ25min	-0.0318	0.016	-2.023	0.043	-0.063	-0.001
direction_same	0.1033	0.013	8.208	0.000	0.079	0.128

```
=====
Omnibus:          78102.225   Durbin-Watson:          1.682
Prob(Omnibus):    0.000   Jarque-Bera (JB):          1293.019
Skew:             -0.228   Prob(JB):          1.68e-281
Kurtosis:         1.504   Cond. No.          24.7
=====
```

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
 [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In the fifth iteration, age is least statistically significant because it has highest p-value of 0.166 which is greater than alpha of 0.05. age should be dropped.

```
x_stepwise_6 = x_stepwise_5.drop("age", axis=1)
result_stepwise_6 = sm.OLS(y_stepwise, x_stepwise_6).fit()
print(result_stepwise_6.summary())
```



OLS Regression Results

```
=====
Dep. Variable:          Y   R-squared (uncentered):          0.587
Model:                  OLS   Adj. R-squared (uncentered):          0.586
Method:                 Least Squares   F-statistic:          1124.
Date:                   Mon, 11 Nov 2024   Prob (F-statistic):          0.00
Time:                   05:49:28   Log-Likelihood:          -8810.4
=====
```

No. Observations:	12684	AIC:	1.765e+04
Df Residuals:	12668	BIC:	1.777e+04
Df Model:	16		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
destination	0.0573	0.007	8.395	0.000	0.044	0.071
passanger	0.0476	0.005	9.511	0.000	0.038	0.057
weather	0.1179	0.007	16.677	0.000	0.104	0.132
temperature	0.0266	0.006	4.178	0.000	0.014	0.039
time	-0.0094	0.004	-2.393	0.017	-0.017	-0.002
coupon	0.0545	0.003	16.551	0.000	0.048	0.061
expiration	-0.1405	0.009	-15.755	0.000	-0.158	-0.123
gender	0.0869	0.009	10.147	0.000	0.070	0.104
maritalStatus	0.0644	0.005	12.772	0.000	0.054	0.074
has_children	0.0496	0.009	5.403	0.000	0.032	0.068
education	0.0205	0.002	9.148	0.000	0.016	0.025
CoffeeHouse	-0.0362	0.003	-12.978	0.000	-0.042	-0.031
RestaurantLessThan20	0.0153	0.004	4.082	0.000	0.008	0.023
toCoupon_GEQ15min	0.0248	0.009	2.636	0.008	0.006	0.043
toCoupon_GEQ25min	-0.0313	0.016	-1.995	0.046	-0.062	-0.001
direction_same	0.1040	0.013	8.269	0.000	0.079	0.129

Omnibus:	78704.924	Durbin-Watson:	1.682
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1288.500
Skew:	-0.228	Prob(JB):	1.61e-280
Kurtosis:	1.507	Cond. No.	22.4

Notes:

[1] R² is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Therefore, age, income, Restaurant20To50, CarryAway and Bar need to be dropped from encoded_in_vehicle_coupon_data to create a dataset that is dimensionally reduced due to stepwise regression.

```
stepwise_reduced_data = encoded_in_vehicle_coupon_data.drop(["Bar", "Restaurant20To50", "CarryAway", "income", "age"], axis=1)
stepwise_reduced_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12684 entries, 0 to 12683
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   destination            12684 non-null  float64
1   passanger              12684 non-null  float64
2   weather                12684 non-null  float64
3   temperature            12684 non-null  float64
4   time                   12684 non-null  float64
5   coupon                 12684 non-null  float64
6   expiration             12684 non-null  float64
7   gender                 12684 non-null  float64
8   maritalStatus          12684 non-null  float64
9   has_children           12684 non-null  float64
10  education              12684 non-null  float64
11  CoffeeHouse            12684 non-null  float64
12  RestaurantLessThan20   12684 non-null  float64
13  toCoupon_GEQ15min      12684 non-null  float64
14  toCoupon_GEQ25min      12684 non-null  float64
15  direction_same         12684 non-null  float64
16  Y                      12684 non-null  float64
dtypes: float64(17)
memory usage: 1.6 MB
```

```
#Now I will do forward feature selection for dimensionality reduction of encoded_in_vehicle_coupon_data
from sklearn.feature_selection import SequentialFeatureSelector as sfs
from sklearn.linear_model import LinearRegression
```

```
lreg = LinearRegression()
X_forward_feature_selection = encoded_in_vehicle_coupon_data.drop(['Y'], axis=1)
y_forward_feature_selection = encoded_in_vehicle_coupon_data['Y']
#I wasn to keep the same number of features as stepwise regression, so I picked 16
sfs1 = sfs(estimator=lreg, n_features_to_select=16, direction='forward', scoring='neg_mean_squared_error')
sfs1 = sfs1.fit(X_forward_feature_selection, y_forward_feature_selection)
selected_features_indices = sfs1.get_support(indices=True)
feature_names = X_forward_feature_selection.columns
# Get the names of the selected features
selected_features_forward = feature_names[selected_features_indices]
print("Selected feature names:", selected_features_forward)
```

```
Selected feature names: Index(['destination', 'passanger', 'weather', 'time', 'coupon', 'expiration',  
                             'gender', 'age', 'has_children', 'education', 'income', 'CoffeeHouse',  
                             'CarryAway', 'Restaurant20To50', 'toCoupon_GEQ15min',  
                             'toCoupon_GEQ25min'],  
                             dtype='object')
```

```
#Below code snippet from https://stackoverflow.com/questions/72985935/how-to-run-machine-learning-algorithms-in-gpu  
unselected_features_forward = list(set(feature_names) - set(selected_features_forward))  
print("Unselected feature names by forward feature selection ", unselected_features_forward)
```

```
Unselected feature names by forward feature selection ['maritalStatus', 'direction_same', 'temperature', 'Bar', 'RestaurantLessThan20']
```

```
forward_feature_reduced_data = encoded_in_vehicle_coupon_data.drop(["Bar", "RestaurantLessThan20", "direction_same", "temperature", "maritalStatus"], axis=1)  
forward_feature_reduced_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 12684 entries, 0 to 12683  
Data columns (total 17 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   destination            12684 non-null  float64  
1   passanger              12684 non-null  float64  
2   weather                12684 non-null  float64  
3   time                   12684 non-null  float64  
4   coupon                 12684 non-null  float64  
5   expiration              12684 non-null  float64  
6   gender                 12684 non-null  float64  
7   age                    12684 non-null  float64  
8   has_children            12684 non-null  float64  
9   education              12684 non-null  float64  
10  income                  12684 non-null  float64  
11  CoffeeHouse             12684 non-null  float64  
12  CarryAway               12684 non-null  float64  
13  Restaurant20To50        12684 non-null  float64  
14  toCoupon_GEQ15min       12684 non-null  float64  
15  toCoupon_GEQ25min       12684 non-null  float64  
16  Y                       12684 non-null  float64  
dtypes: float64(17)  
memory usage: 1.6 MB
```

Stepwise Regression and Forward Feature selection for dimensionality reduction gave different results resulting in different columns being dropped. Let's now see the results of Backward Feature elimination.

```
X_backward_feature_elimination = encoded_in_vehicle_coupon_data.drop(['Y'], axis=1)  
y_backward_feature_elimination = encoded_in_vehicle_coupon_data['Y']  
#I wasn't to keep the same number of features as stepwise regression and forward feature selection for comparison, so I picked 16  
sfs2 = sfs(estimator=lreg, n_features_to_select=16, direction = 'backward', scoring='neg_mean_squared_error')  
sfs2 = sfs2.fit(X_backward_feature_elimination, y_backward_feature_elimination)  
selected_features_indices_2 = sfs2.get_support(indices=True)  
feature_names_2 = X_backward_feature_elimination.columns  
# Get the names of the selected features  
selected_features_backward = feature_names_2[selected_features_indices_2]  
print("Selected feature names for backward feature elimination:", selected_features_backward)
```

```
Selected feature names for backward feature elimination: Index(['destination', 'passanger', 'weather', 'time', 'coupon', 'expiration',  
                             'gender', 'age', 'has_children', 'education', 'income', 'CoffeeHouse',  
                             'CarryAway', 'Restaurant20To50', 'toCoupon_GEQ15min',  
                             'toCoupon_GEQ25min'],  
                             dtype='object')
```

```
#Below code snippet from https://stackoverflow.com/questions/72985935/how-to-run-machine-learning-algorithms-in-gpu  
unselected_features_backward = list(set(feature_names_2) - set(selected_features_backward))  
print("Unselected feature names for backward feature elimination ", unselected_features_backward)
```

```
Unselected feature names for backward feature elimination ['maritalStatus', 'direction_same', 'temperature', 'Bar', 'RestaurantLessThan20']
```

```
backward_feature_reduced_data = encoded_in_vehicle_coupon_data.drop(["Bar", "RestaurantLessThan20", "direction_same", "temperature", "maritalStatus"], axis=1)  
backward_feature_reduced_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 12684 entries, 0 to 12683  
Data columns (total 17 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   destination            12684 non-null  float64  
1   passanger              12684 non-null  float64  
2   weather                12684 non-null  float64  
3   time                   12684 non-null  float64  
4   coupon                 12684 non-null  float64
```

```

5   expiration      12684 non-null   float64
6   gender           12684 non-null   float64
7   age              12684 non-null   float64
8   has_children     12684 non-null   float64
9   education        12684 non-null   float64
10  income           12684 non-null   float64
11  CoffeeHouse      12684 non-null   float64
12  CarryAway        12684 non-null   float64
13  Restaurant20To50 12684 non-null   float64
14  toCoupon_GEQ15min 12684 non-null   float64
15  toCoupon_GEQ25min 12684 non-null   float64
16  Y                12684 non-null   float64
dtypes: float64(17)
memory usage: 1.6 MB

```

It is thus observed that the results of Forward feature selection and Backward feature elimination are the same and give the same reduced dataset.


I will now run the classification algorithms on the 2 dimensionally reduced datasets for comparison. The 2 data sets are: the one reduced using stepwise regression and the other reduced using backward feature elimination/forward feature selection.

Random forest on dataset reduced using stepwise regression

```

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import roc_auc_score, roc_curve
X_random_forest_stepwise_data = stepwise_reduced_data.drop('Y', axis=1)
y_random_forest_stepwise_data = stepwise_reduced_data['Y']
X_random_forest_stepwise_train, X_random_forest_stepwise_test, y_random_forest_stepwise_train, y_random_forest_stepwise_test = train_test_split(X_random_forest_stepwise_data, y_random_forest_stepwise_data, random_state=50)
rf_stepwise = RandomForestClassifier(n_estimators=100, random_state=50)
rf_stepwise.fit(X_random_forest_stepwise_train, y_random_forest_stepwise_train)
y_pred_stepwise = rf_stepwise.predict(X_random_forest_stepwise_test)
print("Accuracy:", accuracy_score(y_random_forest_stepwise_test, y_pred_stepwise))
print("Classification Report:\n", classification_report(y_random_forest_stepwise_test, y_pred_stepwise))
y_pred_prob_random_forest_stepwise = rf_stepwise.predict_proba(X_random_forest_stepwise_test)[:, 1] #make predictions
auc_random_forest_stepwise = roc_auc_score(y_random_forest_stepwise_test, y_pred_prob_random_forest_stepwise)
print("AUC:", auc_random_forest_stepwise)

```



```

Accuracy: 0.6862435947970044
Classification Report:

```

	precision	recall	f1-score	support
0.0	0.64	0.61	0.63	1087
1.0	0.72	0.74	0.73	1450
accuracy			0.69	2537
macro avg	0.68	0.68	0.68	2537
weighted avg	0.68	0.69	0.69	2537


AUC: 0.7351197538305365

Random forest on dataset reduced using forward feature selection/backward feature elimination

```

X_random_forest_backward_data = backward_feature_reduced_data.drop('Y', axis=1)
y_random_forest_backward_data = backward_feature_reduced_data['Y']
X_random_forest_backward_train, X_random_forest_backward_test, y_random_forest_backward_train, y_random_forest_backward_test = train_test_split(X_random_forest_backward_data, y_random_forest_backward_data, random_state=50)
rf_backward = RandomForestClassifier(n_estimators=100, random_state=50)
rf_backward.fit(X_random_forest_backward_train, y_random_forest_backward_train)
y_pred_backward = rf_backward.predict(X_random_forest_backward_test)
print("Accuracy:", accuracy_score(y_random_forest_backward_test, y_pred_backward))
print("Classification Report:\n", classification_report(y_random_forest_backward_test, y_pred_backward))
y_pred_prob_random_forest_backward = rf_backward.predict_proba(X_random_forest_backward_test)[:, 1] #make predictions
auc_random_forest_backward = roc_auc_score(y_random_forest_backward_test, y_pred_prob_random_forest_backward)
print("AUC:", auc_random_forest_backward)

```



```

Accuracy: 0.733543553803705
Classification Report:

```

	precision	recall	f1-score	support
0.0	0.70	0.65	0.68	1087
1.0	0.75	0.79	0.77	1450
accuracy			0.73	2537
macro avg	0.73	0.72	0.73	2537
weighted avg	0.73	0.73	0.73	2537

AUC: 0.7942876629762395

Logistic Regression on dataset reduced using stepwise regression

```
from sklearn.linear_model import LogisticRegression
X_logistic_reg_stepwise_data = stepwise_reduced_data.drop('Y', axis=1)
y_logistic_reg_stepwise_data = stepwise_reduced_data['Y']
X_logistic_reg_stepwise_train, X_logistic_reg_stepwise_test, y_logistic_reg_stepwise_train, y_logistic_reg_stepwise_test = train_test_split(X_logistic_reg_stepwise_data, y_logistic_reg_stepwise_data, test_size=0.2, random_state=42)
logistic_stepwise_model = LogisticRegression()
logistic_stepwise_model.fit(X_logistic_reg_stepwise_train, y_logistic_reg_stepwise_train)
y_pred_logistic_reg_stepwise = logistic_stepwise_model.predict(X_logistic_reg_stepwise_test)
print("Accuracy:", accuracy_score(y_logistic_reg_stepwise_test, y_pred_logistic_reg_stepwise))
print("Classification Report:\n", classification_report(y_logistic_reg_stepwise_test, y_pred_logistic_reg_stepwise))
y_pred_prob_logistic_reg_stepwise = logistic_stepwise_model.predict_proba(X_logistic_reg_stepwise_test)[:, 1] #make predictions
auc_logistic_reg_stepwise = roc_auc_score(y_logistic_reg_stepwise_test, y_pred_prob_logistic_reg_stepwise)
print("AUC:", auc_logistic_reg_stepwise)
```

➤ Accuracy: 0.6243594797004336

Classification Report:				
	precision	recall	f1-score	support
0.0	0.59	0.42	0.49	1087
1.0	0.64	0.77	0.70	1450
accuracy			0.62	2537
macro avg	0.61	0.60	0.60	2537
weighted avg	0.62	0.62	0.61	2537

AUC: 0.66001966817879

Logistic Regression on dataset reduced using forward feature selection/backward feature elimination

```
X_logistic_reg_backward_data = backward_feature_reduced_data.drop('Y', axis=1)
y_logistic_reg_backward_data = backward_feature_reduced_data['Y']
X_logistic_reg_backward_train, X_logistic_reg_backward_test, y_logistic_reg_backward_train, y_logistic_reg_backward_test = train_test_split(X_logistic_reg_backward_data, y_logistic_reg_backward_data, test_size=0.2, random_state=42)
logistic_backward_model = LogisticRegression()
logistic_backward_model.fit(X_logistic_reg_backward_train, y_logistic_reg_backward_train)
y_pred_logistic_reg_backward = logistic_backward_model.predict(X_logistic_reg_backward_test)
print("Accuracy:", accuracy_score(y_logistic_reg_backward_test, y_pred_logistic_reg_backward))
print("Classification Report:\n", classification_report(y_logistic_reg_backward_test, y_pred_logistic_reg_backward))
y_pred_prob_logistic_reg_backward = logistic_backward_model.predict_proba(X_logistic_reg_backward_test)[:, 1] #make predictions
auc_logistic_reg_backward = roc_auc_score(y_logistic_reg_backward_test, y_pred_prob_logistic_reg_backward)
print("AUC:", auc_logistic_reg_backward)
```

➤ Accuracy: 0.6267244777296019

Classification Report:				
	precision	recall	f1-score	support
0.0	0.59	0.42	0.49	1087
1.0	0.64	0.78	0.70	1450
accuracy			0.63	2537
macro avg	0.62	0.60	0.60	2537
weighted avg	0.62	0.63	0.61	2537

AUC: 0.6605948037940551

k Nearest Neighbours (k-NN) on dataset reduced using stepwise regression

```
from sklearn.neighbors import KNeighborsClassifier
X_knn_stepwise_data = stepwise_reduced_data.drop('Y', axis=1)
y_knn_stepwise_data = stepwise_reduced_data['Y']
X_knn_stepwise_train, X_knn_stepwise_test, y_knn_stepwise_train, y_knn_stepwise_test = train_test_split(X_knn_stepwise_data, y_knn_stepwise_data, test_size=0.2, random_state=42)
k = 5 # Number of neighbors
knn_stepwise = KNeighborsClassifier(n_neighbors=k)
knn_stepwise.fit(X_knn_stepwise_train, y_knn_stepwise_train)
y_pred_knn_stepwise = knn_stepwise.predict(X_knn_stepwise_test)
print("Accuracy:", accuracy_score(y_knn_stepwise_test, y_pred_knn_stepwise))
print("Classification Report:\n", classification_report(y_knn_stepwise_test, y_pred_knn_stepwise))
y_pred_prob_knn_stepwise = knn_stepwise.predict_proba(X_knn_stepwise_test)[:, 1] #make predictions
auc_knn_stepwise = roc_auc_score(y_knn_stepwise_test, y_pred_prob_knn_stepwise)
print("AUC:", auc_knn_stepwise)
```

➤ Accuracy: 0.6846669294442255

Classification Report:				
	precision	recall	f1-score	support
0.0	0.64	0.59	0.62	1087
1.0	0.71	0.75	0.73	1450

accuracy			0.68	2537
macro avg	0.68	0.67	0.67	2537
weighted avg	0.68	0.68	0.68	2537

AUC: 0.7027078640992291

k Nearest Neighbours (k-NN) on dataset reduced using forward feature selection/backward feature elimination

```
X_knn_backward_data = backward_feature_reduced_data.drop('Y', axis=1)
y_knn_backward_data = backward_feature_reduced_data['Y']
X_knn_backward_train, X_knn_backward_test, y_knn_backward_train, y_knn_backward_test = train_test_split(X_knn_backward_data, y_knn_backward_data, test_size=0.2, random_state=42)
k = 5 # Number of neighbors
knn_backward = KNeighborsClassifier(n_neighbors=k)
knn_backward.fit(X_knn_backward_train, y_knn_backward_train)
y_pred_knn_backward = knn_backward.predict(X_knn_backward_test)
print("Accuracy:", accuracy_score(y_knn_backward_test, y_pred_knn_backward))
print("Classification Report:\n", classification_report(y_knn_backward_test, y_pred_knn_backward))
y_pred_prob_knn_backward = knn_backward.predict_proba(X_knn_backward_test)[:, 1] #make predictions
auc_knn_backward = roc_auc_score(y_knn_backward_test, y_pred_prob_knn_backward)
print("AUC:", auc_knn_backward)
```

```

➡ Accuracy: 0.6460386283011431
Classification Report:
      precision    recall  f1-score   support

    0.0         0.60      0.54      0.57      1087
    1.0         0.68      0.73      0.70      1450

 accuracy         0.65      0.65      0.65      2537
  macro avg       0.64      0.63      0.63      2537
 weighted avg     0.64      0.65      0.64      2537

AUC: 0.6774247374932589
```

Naïve Bayes on dataset reduced using stepwise regression. I will use Gaussian Naïve Bayes.

```
from sklearn.naive_bayes import GaussianNB
X_nb_stepwise_data = stepwise_reduced_data.drop('Y', axis=1)
y_nb_stepwise_data = stepwise_reduced_data['Y']
X_nb_stepwise_train, X_nb_stepwise_test, y_nb_stepwise_train, y_nb_stepwise_test = train_test_split(X_nb_stepwise_data, y_nb_stepwise_data, test_size=0.2, random_state=42)
nb_stepwise_model = GaussianNB()
nb_stepwise_model.fit(X_nb_stepwise_train, y_nb_stepwise_train)
y_pred_nb_stepwise = nb_stepwise_model.predict(X_nb_stepwise_test)
print("Accuracy:", accuracy_score(y_nb_stepwise_test, y_pred_nb_stepwise))
print("Classification Report:\n", classification_report(y_nb_stepwise_test, y_pred_nb_stepwise))
y_pred_prob_nb_stepwise = nb_stepwise_model.predict_proba(X_nb_stepwise_test)[:, 1] #make predictions
auc_nb_stepwise = roc_auc_score(y_nb_stepwise_test, y_pred_prob_nb_stepwise)
print("AUC:", auc_nb_stepwise)
```

```

➡ Accuracy: 0.5880961765865195
Classification Report:
      precision    recall  f1-score   support

    0.0         0.53      0.34      0.41      1087
    1.0         0.61      0.78      0.68      1450

 accuracy         0.57      0.56      0.57      2537
  macro avg       0.57      0.56      0.55      2537
 weighted avg     0.58      0.59      0.57      2537

AUC: 0.6340354661675602
```

Naïve Bayes on dataset reduced using forward feature selection/backward feature elimination

```
X_nb_backward_data = backward_feature_reduced_data.drop('Y', axis=1)
y_nb_backward_data = backward_feature_reduced_data['Y']
X_nb_backward_train, X_nb_backward_test, y_nb_backward_train, y_nb_backward_test = train_test_split(X_nb_backward_data, y_nb_backward_data, test_size=0.2, random_state=42)
nb_backward_model = GaussianNB()
nb_backward_model.fit(X_nb_backward_train, y_nb_backward_train)
y_pred_nb_backward = nb_backward_model.predict(X_nb_backward_test)
print("Accuracy:", accuracy_score(y_nb_backward_test, y_pred_nb_backward))
print("Classification Report:\n", classification_report(y_nb_backward_test, y_pred_nb_backward))
y_pred_prob_nb_backward = nb_backward_model.predict_proba(X_nb_backward_test)[:, 1] #make predictions
auc_nb_backward = roc_auc_score(y_nb_backward_test, y_pred_prob_nb_backward)
print("AUC:", auc_nb_backward)
```

```

Accuracy: 0.5904611746156878
Classification Report:

```

	precision	recall	f1-score	support
0.0	0.53	0.34	0.42	1087
1.0	0.61	0.78	0.68	1450
accuracy			0.59	2537
macro avg	0.57	0.56	0.55	2537
weighted avg	0.58	0.59	0.57	2537

```

AUC: 0.620200133713615

```

Decision Tree on dataset reduced using stepwise regression

```

from sklearn.tree import DecisionTreeClassifier
X_decision_tree_stepwise_data = stepwise_reduced_data.drop('Y', axis=1)
y_decision_tree_stepwise_data = stepwise_reduced_data['Y']
X_dt_stepwise_train, X_dt_stepwise_test, y_dt_stepwise_train, y_dt_stepwise_test = train_test_split(X_decision_tree_stepwise_data, y_decision_tree_stepwise_data, test_size=0.3, random_state=42)
dt_stepwise_model = DecisionTreeClassifier()
dt_stepwise_model.fit(X_dt_stepwise_train, y_dt_stepwise_train)
y_pred_dt_stepwise = dt_stepwise_model.predict(X_dt_stepwise_test)
print("Accuracy:", accuracy_score(y_dt_stepwise_test, y_pred_dt_stepwise))
print("Classification Report:\n", classification_report(y_dt_stepwise_test, y_pred_dt_stepwise))
y_pred_prob_dt_stepwise = dt_stepwise_model.predict_proba(X_dt_stepwise_test)[:, 1] #make predictions
auc_dt_stepwise = roc_auc_score(y_dt_stepwise_test, y_pred_prob_dt_stepwise)
print("AUC:", auc_dt_stepwise)

```

```

Accuracy: 0.6590461174615688
Classification Report:

```

	precision	recall	f1-score	support
0.0	0.60	0.63	0.61	1087
1.0	0.71	0.68	0.70	1450
accuracy			0.66	2537
macro avg	0.65	0.66	0.65	2537
weighted avg	0.66	0.66	0.66	2537

```

AUC: 0.659372521650858

```

Decision Tree on dataset reduced using forward feature selection/backward feature elimination

```

X_decision_tree_backward_data = backward_feature_reduced_data.drop('Y', axis=1)
y_decision_tree_backward_data = backward_feature_reduced_data['Y']
X_dt_backward_train, X_dt_backward_test, y_dt_backward_train, y_dt_backward_test = train_test_split(X_decision_tree_backward_data, y_decision_tree_backward_data, test_size=0.3, random_state=42)
dt_backward_model = DecisionTreeClassifier()
dt_backward_model.fit(X_dt_backward_train, y_dt_backward_train)
y_pred_dt_backward = dt_backward_model.predict(X_dt_backward_test)
print("Accuracy:", accuracy_score(y_dt_backward_test, y_pred_dt_backward))
print("Classification Report:\n", classification_report(y_dt_backward_test, y_pred_dt_backward))
y_pred_prob_dt_backward = dt_backward_model.predict_proba(X_dt_backward_test)[:, 1] #make predictions
auc_dt_backward = roc_auc_score(y_dt_backward_test, y_pred_prob_dt_backward)
print("AUC:", auc_dt_backward)

```

```

Accuracy: 0.6637761135199054
Classification Report:

```

	precision	recall	f1-score	support
0.0	0.60	0.64	0.62	1087
1.0	0.72	0.68	0.70	1450
accuracy			0.66	2537
macro avg	0.66	0.66	0.66	2537
weighted avg	0.67	0.66	0.66	2537

```

AUC: 0.6611940487897725

```

Random forest on dataset reduced using forward feature selection/backward feature elimination had the highest accuracy and AUC value closest to 1.